



The Memory Rail for Autonomous Agents

Your memory. Your rules.

Abstract

Most enterprise “misses” are not model or dashboard failures; they are **interoperability failures**—functions, systems, and partners can’t understand each other’s decisions in time to act together. This shows up as status-chasing, duplicated effort, stockouts, expedite fees, disputes, and compliance gaps. In the future, filled with autonomous agents the problem will only get worse. Agents will operate with their own understanding and perspective potentially exacerbating interoperability challenges. To make autonomous agents actually interoperate across those seams, you need four things: **verifiable agent identity, shared semantics, complex-event coordination, and incentive-aligned settlement**—all operating on a **sharable, task-specific collective memory**. Intellex solves this by **elevating memory to a first-class asset**: a memory layer that turns what agents know (reputation, experience, current state) into governed **Memory Assets** that can be proved, permitted, used, updated, moved, and revoked. The result is **auditable privacy-preserving collaboration** across stacks and chains—so autonomous agents **interoperate** without sharing raw data.

Intellex has built the memory rails enabling cross-chain agent interoperability and coordination management.

Executive Summary

Enterprises run as **systems-of-systems**—ERP, WMS, CRM, retail media, portals, suppliers—each on different cadences and incentives. Data are incomplete, decisions are siloed, and human judgment injects noise. The cost is staggering: **poor data quality and downstream decision waste** run in the trillions (IBM estimate for U.S. businesses alone ≈ \$3.1T; MIT Sloan analyses place bad-data drag at **15–25% of revenue**), while **poor decision practices** erode around **20% of annual revenue** and **operational inefficiencies** consume **20–30%** in additional losses.

Intellex addresses the root cause: interoperability. Interoperability between systems; interoperability between business functions; interoperability between business supply chains. Each speaks with different voices: often incomprehensible vocabularies and meta-data that leave functions spending costly human time and financial resources to resolve these misinterpretations. In the era of rapid shifts in business operations costs will continue to grow. In the era of autonomous agents operating in an aspiration coordinated fashion the impact will be catastrophic.

Only with the ability to share memories in a protective fashion and to immutably record memory—actions taken, reputations preserved, expertise offered—will there be the level of trust and confidence in letting agents operate autonomously. Only with interoperability layers that enable translation of problem state, measure impact and provide appropriate incentives, and enable cross-vocabulary interactions will autonomous agents be able to function in a coordinated fashion.

The Intellex protocol is the **memory layer** that lets agents **carry verified reputation, expertise, and experience** across organizations and chains and **\$ITLX** tying utility directly to memory activity—“**no memory activity, no \$ITLX flow; more memory activity, more \$ITLX demand.**”

The Interoperability Problem

Enterprises don't stumble because they lack another dashboard; they stumble because they run on many different memories—marketing's campaign calendar, supply's capacity and lead times, sales' commitments, finance's guardrails—each living in its own tools, rhythms, and incentives. In practice that means truth varies by system, handoffs slip because clocks don't match, and small gaps snowball into missed SLAs, stockouts, and rework. The problem is a systems issue—interoperability—not a spreadsheet issue.

The cost of this misalignment is already immense and rising as autonomous agents proliferate: disparate systems, siloed data, and misaligned incentives are a trillion-dollar drag, and as agents take on more workflows, the lack of shared context will only amplify the problem. Attempts to copy data everywhere (ETL/iPaaS) or to centralize it (lakes/warehouses + “AI on top”) create staleness, version skew, and no portable rights, provenance, or revocation at the moment of action—so teams keep arguing “whose number is correct” instead of executing. What’s missing is a fabric that lets many actors share just enough context, at the right time, with evidence, so the organization can move in sync even when data are imperfect and when artificial intelligence enables new paradigms of problem solving through the collective power of agent collaboration.

The waste is systemic and measurable:

- **Bad data → revenue drag.** Poor data quality costs the U.S. economy ~\$3.1T; for most firms, bad data affect **15–25% of revenue** through rework, delays, and misaligned execution.
- **Decision failure → value erosion.** Poor decision practices erode **~20% of annual revenue**; managers report bad decisions are as common as good ones; mid-level judgment errors alone erase **>3% of profits**.
- **Process friction → lost outcomes.** Companies lose **20–30% of revenue** to operational inefficiencies tied to fragmented systems and handoffs.
- **Reality gap in records.** Data entry errors occur ~5% of the time; CRM stores **68% incomplete, 65% missing fields, 25%+ duplicate** records; **contact decay ~20–30%/year** ensures yesterday’s “truth” is stale.
- **Spreadsheet brittleness.** Audits seldom find an error-free model; **0.4–6.9% per-cell** error rates distort forecasts and reconciliations.

Root cause—made concrete: enterprises are **decentralized, partially observable systems** where each function (and each node in the supply chain) optimizes for a different clock, KPI, and risk model. A single, central “brain” struggles when information is late, incomplete, or politically siloed; even worse, local optimizations routinely conflict—e.g., **Marketing** maximizes ROAS and share-of-voice on fixed campaign windows, **Sales** is paid on bookings and sell-in, **Supply/Operations** protects OTIF and capacity with minimal changeovers, and **Finance** caps working capital and cash-conversion cycles. No one is paid to optimize the whole, so small local moves produce outsized upstream/downstream swings—the classic **bullwhip** pattern.

Interoperability Challenges Inside the Enterprise

It is hard to imagine that business enterprises are challenged by interoperability. Below are a number of examples:

- **Marketing vs. Supply.** Marketing is event-driven (campaign calendars, media windows, SOV/ROAS), which rewards launching on time—even if inventory is constrained. Supply runs an S&OE cadence and values stability (longer runs, fewer changeovers). Result: ad spend against out-of-stock SKUs, expedited freight, and chargebacks.
- **Sales/Trade vs. Finance.** Sales maximizes sell-in, quarter-end bookings, slotting/promo commitments; Finance limits working capital and price dilution, preferring clean margins and shorter cash cycles. Result: channel stuffing, later deductions, and dispute cycles that consume weeks.
- **Operations vs. Product/Channel.** Operations prizes OTIF, fixed ship windows, and DC/FC capacity utilization; Product and Channel teams push resets, NPIs, and assortment changes on retail calendars. Result: mismatched planograms, late cutovers, and rework.
- **Analytics/Planning vs. Execution.** Planning assumes clean, current facts; execution lives with missing fields, judgment noise, and asynchronous updates. The document documents **68% incomplete CRM data, ~20–30% annual contact decay**, and never-error-free spreadsheets—so plans drift the moment they meet reality.

Interoperability Challenges Outside the Enterprise

Friction at the supplier ↔ distributor/retailer seam, where incentives structurally diverge, have even greater interoperability challenges. Here are several examples::

- **Sell-in vs. sell-through. Suppliers** optimize factory utilization, case-pack economics, and quarter-end sell-in; **retailers/distributors** optimize on-shelf availability, category margin, and inventory turns. When demand under-delivers, the retailer faces markdowns and deductions; when demand over-delivers, the supplier faces expedites and OTIF penalties.
- **Trade/retail-media vs. capacity/lead times.** Retail media and trade teams lock spend/placements weeks ahead; suppliers need frozen horizons to plan runs and transportation. Mis-sequenced decisions trigger **OOS advertising**,

late re-allocations, and **DC congestion**—a bullwhip realized.

- **Policy & penalties vs. flexibility. Retailers** enforce OTIF, ASN accuracy, and planogram compliance; **suppliers** ask for flexibility during resets, NPIs, or disruptions. Settlement then moves into **deduction/dispute** workflows—classic interoperability failure that the paper argues is fundamentally **coordination**, not “more data,” with dispute cycles consuming significant hidden margins.

Interoperability Challenges are Forgotten Memories

In the emerging era of artificial intelligence we interpret interoperability challenges as memory problems. Every organization runs on **many memories**—not a single one. Product, supply, finance, marketing, sales, support, and partner teams each keep their own histories, heuristics, and constraints. Those local memories are valuable because they reflect reality at the edge, but they’re also **incomplete and out of sync** with other teams’ memories. What enterprises call “interoperability problems” are, at their core, **memory problems**: the facts and rules one group acts on don’t line up with the facts and rules another group trusts, and the differences are invisible until they cause misses, fire drills, and disputes. Intellex names and standardizes these fragments as **Memory Assets** so they can be **proved, permissioned, licensed, used, updated, moved, and revoked**—with provenance and settlement built in.

The same fragmentation exists **outside** the enterprise. Consumers now carry their own **personal memories** in increasingly capable **personal language models (PLMs)**: model weights fine-tuned on their writing and history plus **memory stacks** (RAG) that store what they know and prefer. PLMs are designed to be **owned by the user**, kept private by default, and synchronized continuously so they “stay in step” with the person. Intellex embraces that reality: personal memories should remain personal—but they should be **licenseable** (on the user’s terms) when doing so benefits the user, like sharing a verified preference, receipt, or eligibility proof in exchange for better service or rewards.

Put together, you get the **real interoperability challenge**: aligning **consumer collective memory** (the personal, user-owned side) with **institutional memory** (the enterprise, multi-team side) in a way that is **synchronized, permissioned, and personalized**—without copying raw data either direction. That alignment is exactly what the Intellex architecture is for. We start by turning both sides into **Memory Assets** with explicit rights and receipts; we **federate** enterprise assets into **institutional memory** so permissions and policies travel with the data across teams and partners; and we let consumers **license narrowly scoped claims** (not their raw data) so agents can personalize decisions in the moment while preserving privacy

and revocation. Every step is auditable and **settled in \$ITLX**, so value flows when memory moves—and only when it moves.

Why is this relevant now? Because **interoperability isn't just technical—it's a business crisis** that already costs trillions in friction, and it will worsen as autonomous agents proliferate unless they share a trustworthy memory. Solving **memory interoperability** inside the enterprise (across functions) and across the enterprise boundary (with consumers) is the key to **true personalization** and the ability to **dynamically respond** to shifts in an enterprise's ecosystem of supply and distribution.. This is how Intellex delivers **trusted, beneficial relationships** between enterprises, and between consumers and enterprises: institutional memory gets smarter and more accountable; personal memory stays sovereign; and agents on both sides of a relationship, whether business-to-business, or business-to-consumer, can coordinate in real time with evidence.

Defining the Potential Solution

Large language models are commoditizing; the race to build larger and larger models is occurring as fast as the shift towards smaller, more proprietary models reflective of niche intelligence. This commoditization value is shifting to whose memory agents run on, while enterprises adopt smaller, domain-specific models and demand ways to use intelligence without exposing raw data. Increasingly there is an expectation that entities, whether consumers or enterprises, expect control and privacy by default.

What's missing: a **sharable, verifiable memory** that carries **constraints, commitments, and permissions** across those boundaries so many agents can act locally yet coordinate globally. The memory supports agents in reputation management , expertise coordination, and experience verification. For enterprises, it is **institutional memory** that travels with provenance and the right to revoke; For consumers, it is **collective memory** that aids in engaging with enterprises that want to communicate and service with them.

Intellex meets these challenges with a memory-first architecture that aligns incentives (through settlement and royalties), enforces permissions (with portable rights and revocation), and synchronizes many memories into one coordinated system of work. In Intellex terms, these are **Memory Assets** (e.g., verified availability snapshots, DC/FC constraints, promo calendars, allowed uses) that are **created, proved, permitted, used, updated, moved, and revoked** with receipts—**each event settled in \$ITLX and anchored on NEAR**—so Marketing's windows, Supply's capacities, Sales' commitments, and Finance's guardrails don't collide in the dark. This is the “**memory rail for AI agents**” positioning:

“Your memory. Your rules.”

1. **Define memories as assets.** Teams and users create Memory Assets (summaries, features, model deltas, attested outcomes) with provenance and permissions.
2. **Federate inside the enterprise.** Individual memories compose into **institutional memory**; policies travel with the data so cross-team coordination doesn't leak privacy.
3. **Bridge to consumers.** Consumers' PLM memories remain theirs; they **license claims** (not raw data) to enterprise agents for personalized decisions, with revocation on demand.
4. **Settle and synchronize.** Each create/permit/use/update/move/revoke event emits a receipt and **settles in \$ITLX**; institutional and personal memories **stay in sync** by design.

And this is why the architecture that follows is merited now.

What We Mean by “Memory”

When we say *memory*, we mean the usable record of past perceptions, actions, and outcomes—structured so it can be retrieved at the right moment to guide the next decision. Memory is not a data dump. It is context that can be looked up by cues (“this SKU, this store, this week”), interpreted consistently (“lead-time means the same thing here as it does there”), and governed (“who is allowed to see or use this, and under what terms”). In practice, it is the substrate that lets many agents—human and machine—coordinate without arguing about whose spreadsheet is “truth.”

There are several kinds of memories that show up in business. **Generalized memory** is what an organization has learned to be generally true—shared concepts, taxonomies, policies, elasticities, and rules of thumb. It is what turns raw events into meaning (“this is a promotion lift, not random noise”). **Institutional memory** is the enterprise's lived history: the specific receipts, commitments, reconciliations, and post-mortems that explain “who did what, when, and why.” It reflects the company's constraints, service levels, and risk tolerances. **Personal memory** is the individual's history and preferences—what a buyer prefers, how a planner sequences work, the styles and thresholds a user actually responds to. Personal memories can, with

permission, be composed into **collective memories** that summarize what a group knows or wants without exposing anyone's raw details.

This is why we call it *memory*, not *intelligence*. Intelligence is the ability to choose among actions; memory is the context those choices depend on. Intelligence without memory is brittle—clever models producing confident guesses in a vacuum. Memory without intelligence is inert—facts that never shape behavior. For interoperability, memory is the durable asset: it is ownable, permissionable, revocable, and auditable; it can be synchronized across teams and time; and it gives many intelligences—human judgment, domain algorithms, small task-specific models—a common, governed ground to stand on. Intelligence will keep changing; your memory is what compounds.

Interoperability breaks down because **multiple memories coexist and diverge**.

Inside any enterprise, marketing's memory (campaign calendars, audience definitions) differs from supply's memory (capacity, lead-times), which differs from finance's memory (cash constraints, guardrails). Each is valid locally but incomplete globally. Outside the enterprise, consumers now carry powerful personal memories in their own assistants—what they prefer, where they've granted consent, how they want to be recognized.

Interoperability Challenges Are Memory Constraints

The major interoperability challenge is to **align business to business institutional memories, or business institutional memory to consumer collective memory** in a way that is synchronized, real-time relevant, and privacy-preserving. Use of large language models to power autonomous agents will mean that agent memory will be critical not only for the agent's own behavior but also for coordination and collaboration in multi-agent systems. Businesses will be faced with the challenge of ensuring that vital institutional memories are instilled in agent reasoning.

Consumers will be faced with the challenge of building personal personas—personal language models—that sufficiently represent them in relationship with businesses.

For the business to consumer relationship that means when a consumer's licensed claims ("deliver weekdays before 9am," "substitute brand B for brand A," "do not use location data beyond 24 hours") meet an enterprise's institutional memory of inventory, routes, and policies, the interaction becomes both relevant and compliant. In B2C, collective memories let brands personalize without hoarding personal data: the consumer's assistant can license a narrow, revocable claim ("eligible for student pricing," "prefers size M in this cut"), and the brand matches it against its institutional memory to fulfill, price, and support accordingly. That alignment is the difference

between generic marketing and a *true* personalized relationship: actions that are correct for *this* person now, while respecting both parties' constraints.

This alignment is also the key to enterprise-to-enterprise coordination. In B2B supply chains, each firm's institutional memory encodes its own constraints (dock hours, quality thresholds, safety stock), commitments (OTIF, price protections), and histories (chargebacks, resolutions). Collective memories allow partners to assemble a **task-specific shared memory** around a shipment, a promo, or a service call—just enough context to decide together, with receipts to settle outcomes later. Imagine, a national grocer and a beverage supplier are preparing a four-week summer promotion on a fast-moving SKU. Each side has its own “memory.” The grocer’s institutional memory includes planograms and reset dates, store-level demand baselines, retail-media placements already sold, OTIF/ASN policies, and historical chargeback patterns. The supplier’s institutional memory includes production slots and changeover costs, ATP by plant, DC capacity and labor windows, lane reliability by carrier, and lot-release history. None of that lives in one place, and the two companies don’t want to copy raw data across their firewalls.

With Intellex, each party turns just the relevant slices of what they know into **Memory Assets** for this promotion: the retailer publishes a permissioned *PromoLift* asset (expected lift by SKU/store/week with confidence), a *PlanogramReset* asset (which stores change on what dates), and an *OTIFPolicy* asset (rules and penalties). The supplier publishes *ATPWindow* (what quantities are actually buildable by week), *DCCapacity* (inbound/outbound slots), *QualityLot* (which lots are cleared), and a *SubstitutionPolicy* for can sizes if aluminum becomes tight. Agent identity and shared schemas (store, SKU, week, DC, lane) let both sides’ planning and logistics agents discover these assets and interpret them the same way.

Now the **task-specific collective memory** exists: a live, shared context—only for the stores, SKUs, and weeks in scope—without exposing any party’s raw systems. As weather and ad inventory shift, **complex-event coordination** kicks in. If the media agent confirms a hotter-than-forecast week while the DC capacity asset shows a bottleneck, a coordination rule proposes shifting volume to a nearby FC and pulling forward a Saturday line run. Because the retailer’s *OTIFPolicy* and the supplier’s *SubstitutionPolicy* are part of the same shared memory, both sides’ agents can choose the compliant move (ship the 10-pack where the shelf can absorb it; hold 12-packs for regions without resets).

Every access and decision produces a receipt. When the retailer’s media agent reads *PromoLift* and *ATPWindow* to finalize placements, that use settles a small royalty to the asset owners; when the supplier relies on *PlanogramReset* to prioritize case-pack, that use settles too. At ship time, delivery events and scan data update

the memory; if an assumption proves wrong (a lot fails release), the supplier **revokes** the affected slice, and the retailer's agents stop using it automatically. This **incentive-aligned settlement** means value flows when memory is used, and each party is paid or penalized according to the rules they published—without haggling over “whose spreadsheet” later.

Shared Agentic Memory in Practice

The result is interoperability in practice: identity ensures the right agents are talking; shared semantics keep meanings aligned; complex-event coordination adapts plans as reality changes; and settlement enforces behavior with audit-grade evidence. Both companies make fewer expedites and miss fewer shelf windows, not because they centralized all their data, but because they synchronized the **specific collective memory** needed for this job—and only for as long as it's needed.

Intellex turns generalized, institutional, and personal knowledge into **Memory Assets** that can be created and proved, permitted and licensed, used, updated, moved, and revoked. Agent identity, shared semantics, and complex-event coordination make those assets retrievable and meaningful at the exact moment work happens; incentive-aligned settlement pays contributors and enforces behavior with audit-grade receipts. The outcome is a sharable, task-specific collective memory that lets many agents interoperate—across functions inside the firm, across firms in a supply chain, and across the enterprise–consumer boundary—so decisions are not just intelligent, but grounded in the right memory, at the right time, under the right rules.

When work captures itself and decisions leave receipts, collective memory compounds and the next cycle starts smarter. At the same time, consumers now carry personal memory in their own personal language models (PLMs)—specialized, user-owned models backed by memory stacks (RAG), increasingly run in hybrid/on-device ways for privacy. Real interoperability therefore means aligning memory across the ecosystem—safely, with user control, and without copying raw data either direction.

This is precisely where Intellex fits. Memories are federated into institutional memory so policies and permissions *travel with the data*; businesses can monetize their memory with parts of their supply chain or exploit memory for competitive advantage. Consumers can license narrow, claim-level access from their PLMs on their terms. The result is context-rich, privacy-preserving coordination—inside the firm and across the firm–consumer boundary—where value flows only when memory moves. That's how we turn stateless computation into continuously learning, accountable actors.

Why is “Cross-Chain” a Particular Problem?

In a decentralized artificial intelligence architecture, agents won’t live on one chain—they’ll sense, decide, and settle across many. That immediately inherits all the heterogeneity of today’s ledgers: different consensus and finality models (probabilistic vs deterministic; reorg-prone vs instant-final), incompatible state models (UTXO vs account), divergent block times and clock semantics, distinct fee markets, data-availability guarantees, and wildly different upgrade/governance cadences. Cross-chain work must reconcile these fundamentals before it can even talk about “business logic.” The same event can be confirmed, reordered, or reverted depending on where you look; the same identity can be representable on one chain and non-portable on another; the same transaction can be cheap in one environment and uneconomical in another. Interoperability isn’t just “wire a bridge”—it’s agreeing on what an event *means*, what “final” means, and how to recover when different chains disagree.

Now add probabilistic AI to deterministic ledgers. Agents are stochastic: they sample, prune, and adapt from context. Ledgers are unforgiving: they require deterministic, replayable state transitions. When stochastic agents coordinate across deterministic systems, basic tasks—event capture, replay, and verification—become brittle. Two agents seeing slightly different prompts, seeds, embeddings, or latencies will produce different outputs; yet the settlement layer wants a single, replayable truth. You can attempt to “prove” off-chain inferences with redundancy, witnesses, or ZK attestation, but cross-chain amplifies the seams: different chains have different proof systems, call data limits, and on-chain verification costs. What looks “provable” on one chain is infeasible on another, so even *how* you vouch for an AI action fragments by venue.

Reputation sharing is harder still. Reputation is only useful if it’s resistant to Sybil attacks, bound to a stable identity, and portable enough for counterparties to price risk. Across chains, identity standards, staking mechanics, and slashing/dispute processes diverge. One chain’s “staked, high-reputation actor” can be a zero-stake unknown elsewhere. Time semantics differ (block time, finality delays), so recency-weighted scores decay inconsistently. Ordering and MEV domains differ, enabling cross-chain bribery and selective disclosure that distort reputation signals. Bridging reputation invites oracle risk: who signs, where is the root of trust, and how do you claw back reputational credit after a cross-chain reorg or adjudication reversal? Without a common memory and receipt format, reputation becomes a collection of unverifiable badges instead of a risk-pricing instrument.

Consolidating “expertise” (the durable learning agents accumulate) faces the same fragmentation. Knowledge formed on Chain A lives under its data-availability,

retention, and privacy regimes; moving it to Chain B isn't just a bridge call—it's reinterpreting schemas, units, encodings, and policies under a different gas, storage, and governance model. Cross-chain standards help at the transport layer, but they don't solve *semantic drift*: what does "lead time," "lift," or "inventory" mean in each domain, version, and jurisdiction? Add legal constraints (who may process what, where), and your "global AI" devolves into per-chain islands with brittle adapters. The result is expertise that can't compound: each chain's agents relearn what the others already know because there's no durable, permissioned, task-scoped memory they can all consult and update with receipts.

Common rebuttals—"just use a universal bridge," "LO will unify consensus," "ZK will prove everything," "FHE lets us compute privately everywhere"—miss practical limits. Bridges shift, not remove, trust; LOs unify transport, not semantics or economics; today's zkML can't cheaply attest rich model behavior on-chain; FHE is still orders of magnitude too slow for most real-time workflows. The defensible path is not to pretend chains (or agents) will homogenize, but to **introduce a chain-agnostic memory and settlement layer**: verifiable agent identity, versioned shared semantics, complex-event coordination that tolerates latency and reorgs, and incentive-aligned settlement—*all operating on a federate, sharable, task-specific collective memory*. That memory must be ownable, permissionable, revocable, and portable, with audit-grade receipts that any chain can anchor. Without that, decentralized AI will multiply silos faster than it multiplies intelligence—and cross-chain interoperability will get worse precisely where businesses need it to get better.

Intellex's Answer

Intellex is the **immutable memory layer** for people, enterprises, and the autonomous agents working for them. We turn knowledge into **Memory Assets** that are **owned, permissioned, and portable across chains**. The protocol defines and enforces a **memory lifecycle—create, prove, permit, use, update, move, revoke**—and **settles each event in \$ITLX**. In short: **no memory activity, no \$ITLX flow; more memory activity, more \$ITLX demand**.

Design principles:

The Intellex Protocol is built on several key principles in solving agent interoperability for businesses:

- **Owner control:** "Your memory. Your rules." Permissions and revocation travel with the asset.

- **Evidence without exposure:** we anchor **provenance and receipts** on-chain; raw data can remain local.
- **Cross-chain by default:** NEAR Intents make **Memory Assets** and **\$ITLX** move wherever work happens.

Key Architectures

For more in-depth technical discussions see the Technology White Paper. Key technologies include:

Memory Assets (the unit of value)

- **Definition.** A Memory Asset is a signed, permissioned artifact of **reputation, expertise, experience, or state** (e.g., a summary, feature set, model delta, or attested outcome) created by a person, enterprise, or agent. Ownership and allowed uses are explicit.
- **Create & Prove.** Asset creation writes **provenance** to NEAR; **attesters stake \$ITLX** and are **slashed** if the attestation is proven wrong.

Permissions, Licensing, and Royalties

- **Permit & License.** Asset owners **grant, price, or revoke** access; **requesters pay in \$ITLX**; micro-royalties flow automatically. **If memory doesn't move, \$ITLX doesn't move.**
- **Use.** Inference calls and consultations are **metered; receipts** bind the use to the permission and the asset version—**auditable without revealing raw data**.

Update, Federate, and Govern

- **Update.** Owners (or governed collectives) post **improvements** (e.g., model fine-tunes, feature deltas) with **reward splits** for contributors in \$ITLX.
- **Federate & Govern.** Individual assets compose into **institutional memory** where **policies travel with the asset**—enabling teams and partners to coordinate **without centralizing data**.

Move & Revoke (cross-chain)

- **Move.** Memory Assets **bridge across chains** using NEAR Intents; **permissions, provenance, and royalty rules** remain intact.

- **Revoke.** Revocations propagate with **audit trails**, ensuring time-bounded access and regulatory compliance.

Settlement and Observability

- **Settlement.** Each lifecycle event **settles in \$ITLX**; **NEAR Intents** provides the initial cross-chain liquidity and a live proving ground for **memory-in-motion**.
- **Receipts.** On-chain anchors plus off-chain references provide **forensic reconstruction** of “who knew what, when”—the basis for enterprise audit packs.

How It Works

The following provides a flow:

1. **Create & Prove** — A team, person, or agent mints a **Memory Asset** (e.g., a summary, feature set, model delta) and writes provenance/attestations to NEAR; attesters **stake \$ITLX** and are **slashed** if wrong.
2. **Permit & License** — Owners grant scoped rights (who can use, for what), set price/terms; requesters pay in **\$ITLX**; rights and audit trails are explicit.
3. **Use** — Agents consume the memory to make decisions; **receipts** bind **who used what version under which permission**—evidence without moving raw data.
4. **Update** — Improvements (e.g., model fine-tunes, better features) propagate with **co-author splits**; contributors share royalties automatically.
5. **Move & Revoke** — Memory bridges across chains via NEAR Intents; revocations propagate with proofs.

NEAR + Intellex — why they matter together

NEAR is a high-performance L1 designed around *Nightshade* sharding and *Doomslug* finality to deliver low fees and fast, predictable confirmation—an execution environment that favors many small, frequent writes (think receipts, attestations, and state updates) rather than rare, monolithic transactions. Contracts run in WASM (e.g., Rust/AssemblyScript), accounts are human-readable, and the protocol continues to evolve specifically toward “chain abstraction.” In NEAR’s chain-abstraction stack, **Chain Signatures** let a NEAR account or contract *sign and*

execute transactions on other chains via decentralized MPC, and **NEAR Intents** route a user's or app's high-level request to the right chain(s) and liquidity while hiding the cross-chain complexity from the user.

Aurora is the EVM environment on NEAR: an Ethereum-compatible runtime (Aurora Engine) and bridging stack that brings familiar Solidity tooling and assets into NEAR's performance envelope. Aurora runs as smart contracts on NEAR and now also offers **Virtual Chains**—customizable, EVM-compatible chains that inherit NEAR's throughput while letting apps isolate workloads, governance, and fee markets. Together with Rainbow Bridge-based connectivity, Aurora gives builders a “best of both”: Ethereum compatibility and NEAR-class speed/fees.

NEAR Intents sits on top of that foundation as the cross-chain abstraction layer. It uses NEAR's Chain Signatures and a solver/relayer marketplace to fulfill user or app “intents” across multiple chains—including non-EVM chains—so the *request* (swap, pay, settle, move) is satisfied without the user micromanaging bridges, wallets, or fee tokens. Aurora Labs has been a **core contributor** and has shipped production flows built on Intents (e.g., an abstracted cross-chain DEX and OTC trading), while the NEAR ecosystem continues to expand Intents' multi-chain coverage (BTC, Solana, Ethereum, NEAR, Base, Arbitrum, Ripple, etc.).

How this maps to Intellex

Intellex's problem domain—**interoperability for autonomous agents**—requires four invariants: (1) verifiable agent identity, (2) shared semantics, (3) complex-event coordination (CEP), and (4) **incentive-aligned settlement**, all operating on a **sharable, task-specific collective memory**. NEAR is where Intellex anchors *receipts* (provenance of “who knew what, when”), *policy updates*, and *royalty/penalty settlements* because the fees and finality profile make high-frequency, fine-grained writes economical. **Chain Signatures** let a NEAR-resident Intellex contract execute on other chains when a workflow spans non-NEAR assets (e.g., paying a supplier on Ethereum or updating a claim on Solana) while keeping the audit root on NEAR. **NEAR Intents** abstracts the cross-chain steps: an Intellex “settle usage & royalties” intent can be fulfilled against whichever chain holds the required liquidity, with the user shielded from bridge-hopping and fee-token juggling.

On the **EVM** side, **Aurora** gives Intellex plug-and-play access to Solidity-based enterprise adapters (ERP, WMS, adtech connectors already written for EVM), and **Virtual Chains** offer a way to isolate verticals or large customers (e.g., a retailer consortium) with custom gas policies and governance—while still anchoring receipts and identity to NEAR and using Intents for cross-chain coordination. Rainbow

Bridge-based connectivity keeps ETH-native assets and identities in reach when a workflow needs them.

Why is Intellex Deploying on the NEAR Stack to Solve a Real Problem?

Consider a supplier–retailer promotion. Intellex turns each party's *live context* (availability snapshots, DC/FC constraints, promo calendars, OTIF rules) into **Memory Assets**. As agents read those assets and act (book capacity, adjust media, reprioritize lanes), Intellex writes **receipts** to NEAR—small, frequent, permanent proofs of access and decisions. When usage of a Memory Asset incurs a micro-royalty (e.g., retailer consumes supplier's availability claim) or a penalty/bonus (e.g., late arrival vs. OTIF window), Intellex **settles on NEAR**. If a downstream step must touch non-NEAR liquidity (e.g., a stablecoin payout on Ethereum), Intellex fires a **NEAR Intent** that uses Chain Signatures to complete the external leg while preserving a single audit spine on NEAR. The **effect** is a production-grade, privacy-preserving coordination loop where value flows whenever memory moves—and every flow is a NEAR transaction.

As Intellex rolls out with additional partners and verticals, three compounding streams accrue to NEAR:

1. **Receipt density.** Each “create, prove, permit, use, update, move, revoke” event emits an on-chain receipt. Interoperability at scale is *receipt-heavy* by design; NEAR’s low-fee finality makes this viable at enterprise frequencies.
2. **Settlement throughput.** Micro-royalties, performance bonds, dispute resolutions, and verifier bounties settle continuously. NEAR processes the bulk of these, with Intents selectively fanning out to external chains when needed.
3. **EVM adjacency.** Aurora hosts EVM-native adapters and, where appropriate, **Virtual Chains** for large customers/consortia. Those environments still anchor receipts and identity to NEAR and use Intents for liquidity routing—creating steady cross-traffic and fees in the NEAR/Intents rail.

Technically, this design avoids the usual cross-chain pitfalls. We’re not forcing all data on-chain: NEAR stores **receipts and state commitments**, not bulk payloads, preserving privacy and speed. We’re not betting on a single bridge: **Intents** is a *routing layer*—able to reach multiple chains, including non-EVMs—secured by **Chain Signatures** MPC and continually extended by the NEAR ecosystem. We’re not

stranding EVM tooling: **Aurora** keeps Solidity code and ETH-native assets in play, while NEAR remains the audit/settlement hub. The result is a pragmatic **chain-abstracted** architecture: NEAR as the coordination ledger; Aurora for EVM execution; Intents to reach out, settle, and return proofs—exactly what enterprise-grade, agent-driven interoperability requires.

Bottom line: Intellex converts messy, cross-company coordination into *memory-driven* workflows that constantly emit receipts and settlements. NEAR's economics and chain-abstraction features make it the right base layer; Aurora ensures EVM parity; and **NEAR Intents** eliminates the user friction of going multi-chain. As deployments scale, the everyday business of forecasting, allocating, shipping, advertising, and reconciling becomes a steady stream of NEAR/Intents transactions—bringing both **transaction volume** and **protocol revenue** to the NEAR ecosystem.

Why the Intellex Protocol Solves Interoperability

Intellex converts “what teams and artificial intelligence models learn” into **governed, tradeable Memory Assets** that **agents can actually act on** across tools/partners **with receipts**. That turns cross-stack work into **auditable, low-friction exchanges**—reducing reconciliations, shortening dispute windows, and accelerating “award → result” cycles.

What changes for the business:

- **Less waste:** fewer status chases and re-keys; reduced **20–30% operational drag** driven by handoffs and stale context.
- **Faster, safer decisions:** decisions tied to **permitted memory + receipts**—not gut or stale records—address the **~20% revenue erosion** from poor decision practices.
- **Audit-ready operations:** “who used what, when, under which right?” is reconstructable in minutes—closing compliance gaps noted in SoR workflows.

Who Wins

- **Enterprises & ecosystems:** a **system of work** built on **licensed memory**, not brittle ETL; privacy-preserving collaboration across partners, with revocation.

- **Builders & PLM owners:** monetize expertise as **Memory Assets**; earn **\$ITLX** micro-royalties on permitted use; keep raw data private.
- **NEAR:** real **transaction volume** from memory events (create/permit/use/update/revoke); **\$ITLX** utility tied 1:1 to protocol usage.

Token Utility & Protocol Economics

\$ITLX = the settlement rail for memory in motion.

Intellex turns what people and enterprises know into **Memory Assets**—versioned summaries, features, model deltas, and attested outcomes that can be proved, permitted, used, updated, moved, and revoked. To make that lifecycle trustworthy at scale, you need an instrument that (a) prices scarce resources, (b) binds identity to behavior over time, and (c) clears obligations between many independent actors **without** a central intermediary. That instrument is the token. **\$ITLX** is not a tip jar; it is the **settlement rail for memory in motion**. Every lifecycle event **moves \$ITLX**; when memory **doesn't** move, **\$ITLX doesn't** either. When memory moves, value moves—instantly, predictably, and with a receipt.

Security: from spam resistance to economic finality

Blockchains give us cryptographic security; tokens add **economic security**. In Intellex, **\$ITLX** provides:

- **Sybil resistance and anti-spam:** every write that changes shared state (registering an asset, issuing a permission, posting an attestation, pushing a revocation) costs a small, predictable amount. That makes large-scale junk traffic expensive and gives honest participants a clean “quiet line.”
- **Bonds that make promises credible:** when an agent publishes a claim (“Lot #A123 is cleared,” “ETA ≤ 48h,” “Promo lift = +22%”), an **attester stakes \$ITLX** behind it. If the claim is proven wrong, a portion is **slashed** and redistributed to challengers and affected parties. Claims become *economically honest by default*.
- **DoS containment and QoS:** access to high-throughput interfaces (hot read/write paths) is **rate-limited by stake** and priced dynamically. That keeps the system available during spikes and ensures the most valuable work is prioritized.

- **Economic finality:** receipts plus settlement close the loop. Once a licensed use settles and the challenge window passes, counterparties can rely on it in accounting and planning—the economic equivalent of finality.

Identity, reputation, expertise, experience—secured by stake and receipts

In Intellex, identity is more than a public key. Agents build a **reputation score** that decays over time and is updated by outcomes: successful, timely, policy-conformant uses of memory push it up; disputes, revocations due to fault, and SLA breaches pull it down. **Reputation is meaningful because it's backed by stake and evidence.** \$ITLX bonds make reputations **costly to fake** and **costly to burn**; receipts tie reputation updates to verifiable events.

“Expertise” and “experience” are not vague labels; they’re **Memory Assets**. When a translator publishes a schema converter, a forecaster posts a lift model, or an ops team shares a constraint summary, each is a signed, versioned artifact with a content hash and policy. **\$ITLX aligns the economics:** you pay to *use* expertise (micro-royalties), you stake to *vouch* for expertise (attestations), and you **earn** when your expertise improves outcomes (reward splits on updates). Over time, that flow differentiates deep expertise (that gets used, proves accurate, and earns) from shallow copies (that sit idle, earn nothing, and can’t afford to stake).

Tokenized memory as a control surface for collaboration

The most fragile moments in multi-agent work are **critical events**: creating a new asset, issuing permissions, using memory to make a decision, updating an asset based on outcomes, moving an asset across chains or org boundaries, and revoking a right. Intellex treats each of these as a **token-metered, receipt-emitting step**:

- **Create & Prove.** Mint a Memory Asset and post a proof (hash, provenance, policy). Attesters **stake \$ITLX**; a small creation fee covers global indexing and future revocation propagation.
- **Permit & License.** Grant scoped rights (who, what, where, for how long). Licensees pay **micro-royalties** in \$ITLX; owners can set dynamic pricing (e.g., higher during peak demand).

- **Use.** Reads that influence a decision are **metered**. A receipt links the use to the permission and asset version; royalties are paid from the caller's channel.
- **Update.** When outcomes arrive (sales lift, delivery confirmations), contributors to the winning asset version receive **reward splits** funded by \$ITLX.
- **Move.** Cross-chain migration of an asset (or its rights) incurs a fee that funds relayers/solvers and ensures **revocations** propagate consistently.
- **Revoke.** Owners (or governed collectives) can revoke rights; a fee ensures revocation signals are durable and quickly gossiped; misuse after revocation can trigger **penalties** paid out of the offender's stake.

This **tokenized control surface** means collaboration is not hand-wavy; it is a sequence of priced, auditable actions with incentives to be precise and timely.

Use case 1: Supplier–retailer promotion (B2B)

A beverage supplier and a national grocer stand up a four-week promo. Each publishes Memory Assets: the retailer posts *PromoLift* and *OTIFPolicy*; the supplier posts *ATPWindow*, *DCCapacity*, and *QualityLots*.

- **Staking for truth.** The supplier's quality attester stakes \$ITLX on each lot-clearance claim. A later failure slashes stake; a correct claim earns nothing extra but preserves reputation.
- **Licensing and royalties.** The retailer's media agent **licenses** *ATPWindow* and *QualityLots*; micro-royalties flow per decision that reads them.
- **Updates and splits.** Actual scans and deliveries update the assets. If an independent forecasting team's model improved accuracy by 10%, their **update split** pays out automatically.
- **Revocation and penalties.** If a lot fails release, the supplier revokes that slice; any downstream use after the revocation timestamp is ineligible for payout and may be penalized.

Each step moves \$ITLX and emits receipts anchored on NEAR. The effect: fewer expedites, fewer disputes, and **economically enforced** coordination—without copying raw data across firewalls.

Use case 2: Personal Language Models (B2C) with rights and royalties

A consumer's assistant (their PLM) holds personal preferences and proofs (e.g., *delivery-before-9am, allergy-to-additive-X, size-M-in-brand-Y*).

- The consumer **licenses** a **claim-level** Memory Asset to a retailer for 30 days, revocable at will.
- The retailer pays **per qualified use** (e.g., when the claim actually steers a recommendation or fulfillment decision).
- If the consumer later revokes, any post-revocation use is provable misuse; **penalties** slash the retailer's bond and pay the consumer.

\$ITLX makes personalization granular, paid, and accountable: brands personalize without hoarding raw data, and consumers are compensated for value actually delivered.

Use case 3: Expertise markets for translators and validators

A third-party team ships a high-quality schema translator between two ERPs.

- They **stake \$ITLX** to vouch for fidelity and latency.
- Every time the translator is used in a licensed workflow, they **earn micro-royalties**.
- If a challenger demonstrates systematic errors, part of the stake is **slashed** and awarded to the challenger and affected buyers.

Over time, the market curates itself: robust components accumulate **reputation and earnings**; low-quality ones can't afford the *cost of being wrong*.

How \$ITLX strengthens identity, reputation, and expertise—mechanistically

Before we enumerate the specifics, it helps to frame *why* \$ITLX matters here. Cryptography gives us addresses; \$ITLX turns those addresses into **economically meaningful identities** that carry history and consequences. In Intellex, every interaction with a Memory Asset emits an auditable receipt and, when material, a token flow: claims are **backed by stake**, licensed uses pay **micro-royalties**, and provable faults trigger **slashing** after a challenge window. That coupling of receipts and settlement binds identity to behavior over time, so **reputation prices risk** (not vanity) and **expertise becomes a yield-bearing asset** that compounds with use

instead of sitting as a static artifact. The mechanisms below describe, step by step, how \$ITLX imposes costs on falsity, rewards verifiable contribution, and keeps identities portable yet accountable across teams, partners, and chains.

1. **Identity with economic weight.** A bare address is free to create; a **reputation-bearing identity** is expensive to forge because it carries stake, history, and active licenses.
2. **Reputation that prices risk.** Scores update only on **receipt-backed events** (not likes or unverifiable badges). Higher reputation reduces collateral requirements and boosts QoS; falling reputation raises costs and constrains concurrency.
3. **Expertise that compounds.** Useful assets get used, earn royalties, and attract collaborators; useless assets get ignored and decay. Because **updates** can pay prior contributors (via reward splits), the system encourages improvement rather than forks that reset reputation to zero.

Managing cross-chain reality without centralizing

Intellex anchors receipts and settlement on NEAR for low-latency, low-fee finality, while **NEAR Intents** handles cross-chain execution when a workflow needs external liquidity or state. \$ITLX remains the **unit of account** for memory lifecycle events; when an intent executes on another chain, the settlement still references the NEAR-anchored receipt. This preserves a single, auditable spine while letting memory—and the work it enables—move to where it's needed.

Why this design survives adversarial pressure

- **Collusion resistance.** Large claims demand larger bonds; verifying parties can be chosen at random from a stake-weighted set to reduce capture.
- **Replay and reorg tolerance.** Economic events finalize after a challenge window; if a counter-chain reorg invalidates an external step, the **bonded solver** bears the loss, not the protocol.
- **Censorship resistance.** Multiple relayers/solvers compete to fulfill intents; fee markets keep the system liveness intact even under load.
- **Privacy by construction.** Tokenized memory moves **rights and receipts**, not raw payloads; ZK or TEE attestations can further minimize disclosure while keeping payouts and penalties enforceable.

A token that only “goes up” is useless to enterprises. \$ITLX is designed to **move when memory moves** and to stay still when it doesn’t. That makes costs legible, benefits attributable, and risk priceable. It’s how Intellex aligns thousands of small, interdependent actions—inside firms, across partners, and at the enterprise-consumer edge—into a coherent, **secure**, and **economically honest** system of work.

In summary, \$ITLX utility is **directly aligned** to protocol usage:

- **Create & Prove.** Register assets; **attesters stake**; slashing for bad attestations.
- **Permit & License.** Pay for access; **owners earn micro-royalties**.
- **Use.** Metered inference/consultation; receipts and payouts.
- **Update.** Fund federated rounds; reward improvements.
- **Move & Revoke.** Pay to bridge and to propagate revocations with audit trails.

This design makes **token value legible** to enterprises: **demand for memory activity** → **demand for \$ITLX**.

Conclusion

People, enterprises, and their agents turn what they know into **Memory Assets**—versioned summaries, features, model deltas, or attested outcomes—then **create & prove, permit & license, use, update, move, and revoke** those assets. Every event produces a **receipt** and settles in **\$ITLX** (so usage is auditable and owners can be paid), while **NEAR Intents** keep memory portable across chains without exposing raw data. The result is **evidence-first coordination**: share *claims and rights*, not databases; carry **institutional memory** to the seam where decisions happen; revoke when policies change; and learn from outcomes over time.

Intellex reduces that drag by making **coordination evidence-based**: rather than copying data across tools, **agents exchange Memory Assets** with explicit permissions, verifiable provenance, and **automatic settlement**. That unlocks:

- **Fewer status chases & disputes.** Proof replaces email archaeology; royalties and rights are set in policy, not negotiated ad hoc.

- **Faster cross-team execution.** Agents act with the **context they are allowed to know**, not whatever a brittle integration copied yesterday.
- **Audit-first operations.** Receipts reconstruct causality in minutes—not weeks—supporting compliance and supplier/partner claims.

Enterprises don't fail because they lack more dashboards; they fail because they operate as **systems-of-systems**—many teams and tools on different clocks and incentives—with a way to share just enough context, at the right time, with evidence. The result is a persistent **interoperability problem**: truth varies by system, handoffs stall because cadences don't match, small gaps snowball into missed SLAs, stockouts, and revenue leakage. “Fix-the-data” programs (new fields, MDM, lakes, BI rebuilds) have delivered disappointing returns precisely because they treat a **systems** issue like a **spreadsheet** issue; the weekly reality remains status-chasing, re-keying, and slide-ware decision-making. Interoperability is **the** business problem of the AI era—and **memory** is the missing primitive. What's missing is a **memory layer** that captures activity passively, coordinates actions, and leaves receipts—so the organization can move in sync even when data are imperfect.

Where conventional approaches break, technically:

- **Copy-based integration (ETL/iPaaS).** Replicating data across apps creates version skew and stale joins; it doesn't carry **rights, provenance**, or **revocation**, so usage can't be audited (or cleanly rolled back) at the moment of action. The business keeps arguing “whose numbers” instead of executing.
- **Central warehouses/lakes + “AI on top.”** Batch pipelines are great for hindsight, not live coordination. They assume humans will keep forms current and “someone else” will interpret tables into actions; in practice, facts drift and execution lags.
- **Event buses without memory.** Pub/sub moves messages but not **institutional memory** (constraints, elasticities, prior outcomes) or **permissions**. Without durable, portable memory, agents re-learn the same lessons and over-share or under-share context.
- **Monolithic “single-brain” agent platforms.** As models commoditize and enterprises adopt **smaller, domain-specific models**, value shifts to **whose memory the agents run on**. A central brain can't respect local policies or privacy at scale. We need owned memory, not another silo.

It is naïve to think “all-on-chain” is a solution. Writing raw data on-chain leaks privacy, raises cost/latency, and fragments across chains. The right pattern is **hybrid**: low-latency off-chain memory and **on-chain receipts** for provenance and settlement, portable via **NEAR Intents**.

Intellex turns what teams and PLMs **know** into **owned, portable, permissioned Memory Assets**—with **\$ITLX** as the settlement rail and **NEAR** as the anchor—so autonomous agents can finally **work together** across stacks and chains. The result: **less waste, faster execution, and audit-ready trust**—under **your rules**.

Use Cases

Here are a number of use cases and how Intellex solves the interoperability challenges

CPG & Retail: Inventory-Aware Media and Promo

- **Problem.** Campaigns launch on stale inventory views; DCs run hot; stockouts rise while spend keeps running. (A textbook “system-of-systems” failure: marketing, supply, retail portals on different clocks.)
- **Intellex. Inventory, logistics, and retail-media agents** consume licensed Memory Assets (e.g., verified availability, DC/FC constraints, expected lift) and **settle** access via \$ITLX; **royalties** flow to the data owners. **Revocations** halt use if permissions change.
- **Outcomes to track.** Fewer stockouts and expedites, shorter dispute windows, higher margin on promo windows—precisely where interoperability waste used to bite.

Healthcare: OR Scheduling & Supplies

- **Problem.** Partial views across scheduling, staffing, and inventory create cancellations and idle OR time.
- **Intellex.** Hospital agents license Memory Assets (block time history, sterile pack readiness, supplier ETA confidence) under strict permissions; usage and improvements settle in \$ITLX; **audit packs** provide evidence of compliance.

Logistics & Field Ops

- **Problem.** Carrier ETAs, yard capacity, and service windows live in different tools; noisy updates cause missed OTIF.
- **Intellex.** Agents license **status and constraint memory; revocations** propagate; **royalties** compensate contributors (e.g., 3PLs) for timely, high-quality signals.

Personal AI / PLMs

- **Problem.** PLMs need **rights-managed sharing** of derived knowledge (summaries, features, model deltas) without exposing raw personal data.
- **Intellex.** Personal assets can be **licensed** to teams or communities with **granular policy, micro-royalties**, and **revoke-on-demand**—all settled in \$ITLX.