

Intellex **Lite Paper**



*A Memory-First Interoperability Protocol
with Usage-Aligned Tokenomics*

This Lite Paper is for informational purposes only and does not constitute an offer to sell, or a solicitation of an offer to buy, any securities, financial instruments, or any other assets, including digital assets or tokens. The information contained herein has not been approved or disapproved by any regulatory authority. No regulatory authority has passed upon or endorsed the merits of this White Paper, the Intellex Protocol, or the ITLX Token.

This Lite Paper is **not intended to be legal, financial, accounting, tax, or investment advice**. Prospective participants are strongly advised to consult with their own legal, financial, accounting, tax, and other professional advisors before making any decision to participate in or interact with the Intellex Protocol or acquire any ITLX Tokens.

Forward-Looking Statements: This White Paper contains forward-looking statements that are based on our current expectations, assumptions, and projections about future events. These statements are subject to significant risks, uncertainties, and assumptions that could cause actual outcomes to differ materially from those expressed or implied in the forward-looking statements. Such risks and uncertainties include, but are not limited to, changes in technology, market conditions, regulatory developments, and the operational and financial performance of the Intellex Protocol. We undertake no obligation to update or revise any forward-looking statements.

Risks: The Intellex Protocol and the ITLX Token are subject to numerous risks, including but not limited to: technological risks, market risks (including price volatility), regulatory risks, security risks, and operational risks. The ITLX Token is not a currency or a security in every jurisdiction and may not be regulated by any financial authority in your country. The value of the ITLX Token can go down as well as up. You could lose all or a substantial portion of your investment.

No Representations or Warranties: The information in this White Paper is provided "as is" and without representations or warranties of any kind, either express or implied, including, but not limited to, implied warranties of merchantability, fitness for a particular purpose, or non-infringement. While we have made every effort to ensure the accuracy and completeness of the information, we make no guarantees, and you should not rely on the information provided herein. The Intellex Protocol and its components are under continuous development and may change significantly.

Jurisdictional Restrictions: The distribution of this White Paper and the potential offering or sale of ITLX Tokens may be restricted by law in certain jurisdictions. It is your responsibility to determine the legality of your participation in the Intellex Protocol and acquisition of ITLX Tokens under the laws of your jurisdiction.

By reading this White Paper, you acknowledge and agree to the foregoing.

Executive Overview

Enterprises don't fail for lack of more dashboards. They fail because they operate as systems-of-systems—many teams, tools, and partners running on different clocks and incentives—without a way to share just enough context, at the right time, with evidence. As autonomous agents spread across those seams, the missing primitive is memory: a sharable, permissioned, auditable working context that lets agents coordinate without copying raw data or arguing about “whose number.”

Intellex is a memory-first interoperability protocol. It turns what people and enterprises know into Memory Assets—signed, permissioned artifacts (summaries, features, verified claims, translators, model deltas) that can be created & proved, permitted & licensed, used, updated, moved, and revoked. Each step emits an audit-grade receipt and, when material, settles in the Intellex token, \$ITLX.

To align incentives end-to-end, \$ITLX flows through three pools that mirror how Intellex is used in practice:

1. **Access Pool (AP):** rewards when enterprises access and settle against institutional memory (licensed reads, adjudications, revocations, cross-agent settlement).
2. **Contribution Rewards Pool (CRP):** rewards individuals and panels when their zero-party claims are actually used (personal/collective memory that powers PLMs and personalization).
3. **Interop Training & Translation Pool (ITP):** rewards builders and validators of translators, schemas, event/reputation/expertise detectors, and cross-agent messaging components (the machinery that solves interoperability).

The result is a protocol where value only moves when memory moves, quality is stake-backed, and every decision has a receipt.

1. Interoperability as a memory problem

1.1 What we mean by "memory"

Memory is the usable record of past perceptions, actions, and outcomes—structured so agents can retrieve it by cues, interpret it consistently, and apply it under policy. Memory is not a data dump. It's a governed context that travels with permission and provenance.

- **Generalized memory:** shared concepts, taxonomies, policies, elasticities—what turns raw events into meaning.
- **Institutional memory:** the enterprise's lived history (receipts, constraints, commitments, reconciliations, post-mortems).
- **Personal memory:** a user's preferences and history (often held by a personal assistant/PLM). Personal memories can compose into collective memories—group summaries that avoid exposing anyone's raw details.

How to Navigate This Document:

Interoperability fails when many valid memories coexist but cannot synchronize at the moment of action. True personalization and cross-company coordination require aligning collective personal memory and institutional memory—safely, with user control, and without copying raw data either way.

1.2 Why current approaches fall short

- Copy-based integration (ETL/iPaaS) → version skew, no portable rights or revocation.
- Central warehouses + "AI on top" → great for hindsight, poor for live coordination.
- Event buses without memory → messages move, institutional memory does not.
- Naive "all-on-chain" → privacy leakage, cost/latency, and fragmented semantics.

Intellex answers with a hybrid, memory-first architecture: low-latency off-chain working memory; on-chain receipts for provenance and settlement; identities and shared semantics to keep meanings aligned; and tokenized incentives to reward what's useful and punish what's not.

2) Architecture overview

2.1 Primitives

- **Identity & Registry** — agent DIDs, capability vectors, attestations, and discovery.
- **Shared Semantics** — versioned schemas and translators to prevent semantic drift.
- **Working Memory & CEP** — a low-latency memory plane plus complex-event rules that trigger the right task at the right time.
- **Receipts & Settlement** — append-only receipts for every material action; settlement flows in \$ITLX.
- **Governance** — parameters, slashing rules, emission weights, and fee routing.

2.2 Memory Asset lifecycle

1. **Create & Prove** — mint an asset (hash, schema, policy); optional bonded attestation.
2. **Permit & License** — grant scoped rights (who, what, where, purpose, duration).
3. **Use** — a licensed read that changes a decision emits a receipt; royalties settle.
4. **Update** — outcomes (e.g., deliveries, lift) improve assets; contributors earn splits.
5. **Move** — bridge rights/receipts across chains; keep provenance intact.
6. **Revoke** — stop future use; propagates instantly and is provably enforced.

2.3 Receipts (the evidence layer)

A **receipt** binds: (who) used (what asset version) under (policy) at (time) to change (decision), plus pointers to prior receipts and proofs. Receipts power audit, payouts, reputation updates, and disputes.

3) The role of the token

3.1 Why a token is necessary

- Neutral settlement across many actors and chains, without a central intermediary.
- Economic security: make spam/abuse expensive; make truth profitable and falsity costly.
- Portable incentives: pay contributors and curators across org boundaries on every qualified use.
- Governance & sustainability: finance public-good components that everyone depends on
- Provenance & trust: clear immutable record of memory, its use and the reputation and expertise of the agents that use it.

3.2 Security properties provided by \$ITLX

- **Sybil resistance:** writes that change of shared state cost money; attacking the protocol is not free.
- **Bonded claims:** attesters stake to vouch for quality (accuracy, latency, conformance); failures are slashed.
- **DoS containment & QoS:** stake-weighted rate limits prioritize valuable work under load.
- **Economic finality:** once the challenge window closes, receipts are economically settled.

3.3 Identity, reputation, expertise, experience

- **Identity with weight:** an address becomes an identity when it carries stake, history, and active licenses.
- **Reputation that prices risk:** updates only on receipt-backed events; good behavior lowers costs and unlocks throughput.
- **Expertise as an asset:** useful translators, detectors, and models earn per use; shallow copies starve.

4) Tokenomics with three pools

4.1 Total supply

1,000,000,000 \$ITLX (fixed at genesis).

4.2 Allocation (high-level)

- **Access Pool (AP)** — 9% reserved for co-emissions against enterprise access receipts.
- **Contribution Rewards Pool (CRP)** — 3% reserved for zero-party claim usage.
- **Interop Training & Translation Pool (ITP)** — 3% reserved for translators/detectors per invocation and bounties.
- **Community, Ecosystem/Grants, Treasury, Team, Investors, Public Sale, Strategic Rebates** — remaining 85% across long-term buckets (vesting, milestones, governance-gated disbursements).

(Percentages and cliffs are shown here to illustrate structure.)

4.3 Fee routing (enterprise pays → who gets paid)

Every enterprise event (licensed read, adjudication, revocation, settlement) pays a usage fee. By default:

- **60%** → participants via AP (royalties to asset owners; attestation fees to bonded attesters; solver fees).
- **10%** → ITP (when interop components were invoked).
- **10%** → CRP (when zero-party claims contributed).
- **15%** → Treasury (ops, risk, insurance).
- **5%** → Burn (issuance discipline).

4.4 Emission mechanics (monthly epochs)

Each pool has a monthly budget—**AP_t, CRP_t, ITP_t**. Emissions are receipts-gated (only released for proven activity), challengeable (with slashing/clawback), and decaying (unused budget partially expires).

General formula for pool X ∈ {AP, CRP, ITP}:

$$\text{Reward}_{i,X,t} = \text{Score}_{i,X,t} \sum_j \text{Score}_{j,X,t} \cdot \min(X_t, \alpha_X \cdot \text{Fees}_{X,t})$$

$\backslash;=\\$

$$\frac{\text{Score}_{i,X,t}}{\sum_j \text{Score}_{j,X,t}} \cdot \min(\text{X}_t, \alpha_X \cdot \text{Fees}_{X,t})$$

- Score = net verified contribution (e.g., qualified reads, invocations, benchmark points).
- X_t = monthly emission cap for the pool.
- $\text{Fees}_{X,t}$ = pool-relevant fee inflows this month.
- α_X = optional matching coefficient (co-emission that scales with fees).

Decay rule:

$$X_{t+1} = \delta_X \cdot \max(0, X_t - E_{X,t}) + \text{NewCap}_{X,t+1} \cdot \delta_X \cdot \max(0, X_t - E_{X,t}) + \text{NewCap}_{X,t+1}$$

- $E_{X,t}$ = emissions actually paid from pool X in month t.
- $\delta_X \in [0,1]$ = rollover factor (e.g., 0.5 → half of the unused budget rolls, half expires).

Challenge window: receipts are pending for τ days. Successful challenges claw back emissions and slash misbehaving stake.

5) How the pools operate in practice

5.1 Access Pool (AP)

- **What it rewards:** when enterprises use institutional memory—licensed reads that change a decision, dispute adjudications, settlement completions, revocation propagations.
- **Who gets paid:** Memory Asset owners (royalties), bonded attesters (attestation fees), solvers/relayers for cross-chain steps, and the protocol (treasury/burn).
- **Example (B2B promo):** A retailer licenses the supplier's ATPWindow and QualityLots assets to plan a 4-week campaign. Each time those assets actually steer a placement or replenishment decision, a qualified receipt is emitted and AP rewards flow: royalties to the supplier, an attestation fee to the validator who staked on QualityLots fidelity, and minor solver fees for cross-chain settlement.

5.2 Contribution Rewards Pool (CRP)

- **What it rewards:** qualified uses of zero-party claims (student eligibility, eco packaging preference, allergy constraint, delivery window).
- **Who gets paid:** the individual (or their assistant) who owns the claim; any attester/validator who bonded quality; the protocol (treasury/burn).
- **Example (B2C personalization):** A consumer licenses "eligible for student pricing until Sep 30" to a brand only for checkout. Each successful discount application logs a receipt and pays a micro-royalty via CRP. Revocation stops future use; misuse after revocation is penalized automatically.

5.3 Interop Training & Translation Pool (ITP)

- **What it rewards:** per-invocation use of translators/ontology mappers, event/reputation/expertise detectors, and cross-agent messaging components; plus bounties for gaps and milestones on fidelity/latency.
- **Who gets paid:** builders and validators who keep interop components accurate and fast under stake.
- **Example (schema translator):** A third-party publishes an ERP-A↔ERP-B attribute mapper, stakes on ≥98.5% fidelity, and charges pennies per invocation. Every production call routes an ITP reward; failing audits slash stake and claw back prior emissions.

6) End-to-end flows (with receipts and settlement)

6.1 Supplier–retailer promotion (B2B)

1. **Publish:** Retailer posts PromoLift, PlanogramReset, OTIFPolicy; Supplier posts ATPWindow, DCCapacity, QualityLots.
2. **Attest:** A bonded validator stakes on QualityLots fidelity (window + SLO).
3. **License & Use:** Retailer's media and supply agents license assets and take decisions; each qualifying use emits receipts.
4. **Settle:** AP fees flow; AP emissions co-match; attesters earn; treasury accrues; a small burn occurs.
5. **Update:** Scan data and shipment events update assets; contributors earn update splits.
6. **Revoke:** If a lot fails release, supplier revokes that slice; further use is blocked and provably penalized.

6.2 PLM-powered checkout (B2C)

1. **Claim:** Consumer's assistant creates "student eligibility" with time-bound proof.
2. **License:** Brand requests checkout-only use; consumer approves.
3. **Use:** Discount applied → receipt logged → CRP micro-royalty.
4. **Revoke:** Permission ends early; any subsequent attempt fails policy checks and is penalizable.

6.3 Translator in the loop (interop)

1. **Deploy:** Builder publishes ERP translator; stakes for fidelity/latency.
2. **Invoke:** Each cross-system read uses a translator → receipt emitted.
3. **Reward:** ITP pays per invocation; poor audits trigger slashing/clawback.

7) Risk model and adversarial resilience

- **Sybil & spam:** writes cost; identities that matter must carry stake and history.
- **Collusion:** random selection of verifiers; cross-checks; escalating bonds for high-impact claims.
- **Replay/reorgs:** economic finality after challenge window; cross-chain steps insured by bonded solvers.
- **Data leakage:** the protocol moves rights and receipts, not raw payloads; ZK/TEE attestations can minimize disclosure.
- **Drift & decay:** time-boxed attestations; rolling audits; automatic decay of emissions budgets to keep incentives current.

8) Governance knobs

- Fee weights (AP/CRP/ITP/treasury/burn).
- Pool budgets & matching coefficients (AP_t , CRP_t , ITP_t and α_X).
- Challenge window τ and slashing ratios by claim class.
- Rollover factor δ (decay).
- Attestation classes (accuracy thresholds, latency SLOs).
- Public-good funding for schemas, benchmarks, reference connectors.

Governance cannot mint around receipts; all emissions are receipts-first.

9) Unit economics (illustrative)

Assume an enterprise runs 20M qualified reads/month across planning, fulfillment, and customer touchpoints.

- Fee: \$0.001/read average \rightarrow \$20,000 AP fees.
 - \$12,000 (60%) to participants (royalties, attesters, solvers)
 - \$2,000 (10%) to ITP components used
 - \$2,000 (10%) to CRP when claims contributed
 - \$3,000 (15%) treasury
 - \$1,000 (5%) burn
- Emissions: AP_t co-matches up to 20% of AP fees (if available), adding \$4,000 worth of \$ITLX to participants.
- Outcome: predictable opex for enterprises; yield for contributors scales with actual usage.

10) Why this will drive real-world volume

- **For enterprises:** fewer expedites/disputes; faster, evidence-first decisions; pay only when memory is used.
- For contributors: ongoing earnings for useful memory; bonded markets where being right pays and being wrong costs.
- For builders: sustainable public-good funding for the interop machinery everyone needs.
- For the ecosystem: every qualified decision emits receipts and settlements—steady on-chain activity linked to real business value.

11) Frequently asked questions

Q: Why not just keep paying vendors to build ETL?

A: ETL copies data but not rights, provenance, or revocation. Intellex pays for governed context that changes a decision—and it leaves a receipt.

Q: Why do individuals get paid?

A: Because zero-party claims (preferences, eligibility, constraints) are often the missing context in personalization and fulfillment. Paying per qualified use rewards value, not data hoarding.

Q: What if claims are wrong or drift?

A: Attesters stake; challenge windows allow clawback and slashing. Revocation stops future use immediately.

Q: Isn't a token overkill?

A: Not when you need neutral settlement, portable incentives, stake-backed quality, and cross-org payments that follow receipts. A token is the minimal instrument that checks all of those boxes.

12) Conclusion

Interoperability isn't a feature—it's a **memory** discipline. Intellex makes memory ownable, permissionable, revocable, and **payable per use**. The \$ITLX token ties that memory to identity, reputation, and settlement so agents can coordinate inside firms, across partners, and with consumers—**with receipts**. Solve memory interoperability and you get what enterprises and users both want: decisions that are correct for this context, now, under clear rules, and a business model where everyone who made that possible gets paid.

14. References

1. "A blockchain-based federated learning mechanism for privacy ...": A study on blockchain-enabled federated learning in healthcare, e.g. diabetes prediction.
<https://www.sciencedirect.com/science/article/abs/pii/S0010482523010958>
2. "Privacy-preserving federated learning for collaborative medical data...": Advances in medical image classification with federated learning.
<https://www.nature.com/articles/s41598-025-97565-4>
3. "AI-enabled blockchain-based federated learning in edge healthcare": A recent arXiv paper detailing blockchain-enhanced federated edge learning.
<https://arxiv.org/html/2506.00416v1>
4. "AI Agents Meet Blockchain: A Survey on Secure and Scalable Collaboration...": MDPI survey covering multi-agent interoperability. <https://www.mdpi.com/1999-5903/17/2/57>
5. "The future of AI will be governed by protocols no one has agreed on yet": Business Insider article on emerging agent communication standards.
<https://www.businessinsider.com/ai-protocol-rules-future-2025-6>
6. "A Survey of AI Agent Protocols": Recent arXiv work on Agent-to-Agent protocols.
<https://arxiv.org/html/2504.16736v2>
7. "Enhancing Blockchain with Layer 2 State Channels Guide": Overview of off-chain scaling via state channels.
<https://antematter.io/blogs/layer-2-state-channels-blockchain-guide>
8. "State Channels: Still Beat All Other Layer-2 Scalability Solutions": Medium article on throughput benefits.
<https://medium.com/yellow-blog/state-channels-still-beat-all-other-layer-2-scalability-solutions-42533d675e92>
9. "Layer-2 Scaling Solutions": Comprehensive guide to rollups, sidechains, and state channels.
<https://transak.com/blog/what-are-layer-2-blockchains-and-how-do-they-work-l1-vs-l2-networks-explained>
10. "AI in Programmatic Advertising: Here's What You Need To Know": Quantcast blog on AI ROI in ad tech. <https://www.quantcast.com/resources/blog/ai-in-programmatic-advertising>

11. "Blockchain Technology In Digital Advertising: Transparency, Fraud Prevention And Trust": ResearchGate article on ad transparency via blockchain.
https://www.researchgate.net/publication/379925338_Blockchain_Technology_In_Digital_Advertising_Transparency_Fraud_Prevention_And_Trust
12. "How Blockchain Can Solve 5 Ad Frauds in Programmatic Ecosystems": Aroscop article on fraud reduction.
<https://www.aroscop.com/how-blockchain-can-solve-5-ad-frauds-in-programmatic-ecosystems>
13. "AI Agents With Crypto Wallets Now Transforming Company Structures": Forbes article on wallet-enabled agents.
<https://www.forbes.com/sites/digital-assets/2025/04/21/ai-agents-with-crypto-wallets-now-transforming-company-structures/>
14. "AI Agents Economy: Why Crypto May Hold The Key To Fund Management": Forbes on AI-driven wallets in finance.
<https://www.forbes.com/sites/tomerniv/2024/11/07/ai-agents-economy-why-crypto-may-hold-the-key-to-fund-management/>
15. "AI agents trading crypto is a hot narrative...": CoinTelegraph feature on memecoin-powered agents.
<https://cointelegraph.com/magazine/hype-ai-agents-crypto-just-turned-into-reality>
16. "The future of the Agent Economy..." (Business Insider) and estimates of trillion-dollar AI automation opportunity.
17. "The rise of privacy preserving AI": Axios overview on federated learning impact.
<https://wwwaxios.com/2019/05/24/federated-learning-ai-privacy>